Final paper

Are Catastrophe Bonds Really Zero-Beta Assets?

Diversification Benefits in Normal Markets and Financial Crises

Submitted to: Mohamed Mnasri
FINA60218A

By:
Tran Nguyen Quoc-Long

December 2026

1. Introduction

A catastrophe (CAT) bond lets an insurer, a reinsurer or a government pass part of its exposure to natural disasters to capital-market investors. The sponsor pays a premium, the investors put up principal that is held as collateral, and if a specified hurricane, earthquake or other event occurs before maturity, part or all of that principal goes to the sponsor (Cummins, 2008). Froot (2001) showed that catastrophe reinsurance in the 1990s was priced far above expected losses and traced the gap to frictions in the supply of reinsurance capital. Securitisation opened catastrophe risk to investors with far more capital, and the market grew accordingly: outstanding CAT bond risk capital rose from about $2.2 billion in 2002 to $23 billion in June 2017 (Drobetz et al., 2020).

To attract those investors, the market offered a simple argument. Whether a hurricane hits Florida or an earthquake hits Tokyo does not depend on the stock market, so the losses on a CAT bond should be unrelated to the losses on stocks and corporate bonds. In CAPM terms, CAT bonds would be zero-beta assets, and any spread above expected losses would reward a risk that adds nothing to the systematic risk of a diversified portfolio. The argument concerns the natural event, while the investor holds a financial security: a contract with a special purpose vehicle, whose principal sits in a collateral account, whose price comes from dealer quotes, and which is held by funds that may need cash at the same moment as everyone else. Each of these features can carry financial stress into the bond when no catastrophe has happened. Whether they do so is an empirical question, and the published answers contradict each other.

Three positions can be distinguished. Haffar and Le Fur (2022) come closest to the zero-beta view, although they do not test it formally: in their sample, CAT bonds add the least to the probability of extreme joint movements with an equity portfolio. Carayannopoulos and Perez (2015) and Gürtler, Hibbeln and Winkelvos (2016) reject it: CAT bond returns became correlated with the stock market after the collapse of Lehman Brothers, and CAT bond premiums move with corporate bond spreads. Drobetz, Schröder and Tegtmeier (2020) also conclude that CAT bonds do not behave like a zero-beta investment. They find, however, that CAT bonds acted as a safe haven against extreme stock market declines after the financial crisis, a result that Feng and Zeng (2025) extend to the COVID-19 period. In this third position the answer depends on the period.

The difference matters for investors and for risk managers. A diversifier lowers the variance of a portfolio on average; a safe haven does not move with the rest of the portfolio in a crash. An investor who buys CAT bonds for the second property when they only have the first is exposed at the worst moment. Risk measurement has the same weakness. Haffar and Le Fur find that their CAT bond fund lowers the 99 percent value-at-risk of an equity portfolio, but if the correlation of CAT bonds with stocks rises when stocks fall, as it did after Lehman’s failure, a value-at-risk estimated on calm data understates the risk of the portfolio in the period when that risk is realised.

The three articles cannot settle the question as they stand, because they do not measure the same object. They use different series (two Swiss Re indices and one fund), frequencies (weekly and

monthly), windows (2002-2013, 2002-2018 and 2011-2020) and definitions of stress (NBER recession dates, extreme weeks of the other asset, and co-kurtosis). None of them separates a crisis in which the collateral of CAT bonds is at risk from one in which investors only need cash. Their conclusions contradict each other, but their estimates need not.

My research question is: **to what extent do CAT bonds retain their diversification benefits during financial stress, and why do published studies reach different answers?** I argue that the disagreement has little to do with whether catastrophes occur and much to do with how catastrophe risk is collateralised, priced and measured. My contribution is a decomposition of the CAT bond return into an insurance core, whose losses are nearly independent of markets for natural perils, and a financial wrapper, whose market exposure depends on the collateral structure and on the type of stress. The decomposition sorts the samples of the three articles into four regimes, yields a prediction that separates it from the view that 2008 was a one-off, and leads to proposals for investors. The answer it gives is conditional: CAT bonds remain a reliable diversifier, but their safe-haven property requires safe collateral and weakens when holders are forced to sell.

Section 2 explains how CAT bonds work, defines the properties that the literature tests and presents the three articles. Section 3 assesses the articles, develops the decomposition, applies it to the evidence and proposes tests and solutions. Section 4 concludes.

2. Literature review

2.1 How a CAT bond works, and where financial risk can enter

A sponsor sets up a special purpose vehicle (SPV) that sells notes to investors and signs a reinsurance contract with the sponsor. The proceeds of the notes are placed in a collateral trust account. Investors receive a floating coupon, a money-market rate plus a risk premium, and the bonds usually mature in three to four years (Drobetz et al., 2020). If a qualifying event occurs, the collateral is released to the sponsor and investors lose part or all of their principal. The trigger can be the sponsor's own losses, an industry loss index, a modelled loss or a physical parameter such as wind speed (Cummins, 2008; Cummins and Weiss, 2009).

Several parts of this structure have nothing to do with nature. One is the **collateral**. Before 2009, the vast majority of CAT bonds used a total return swap: an investment bank converted the return on the collateral into a money-market rate and was supposed to absorb changes in the value of the collateral. The collateral itself often had much longer maturities than the bond and could include highly rated asset-backed and mortgage-backed securities, with little monitoring and no obligation for the swap counterparty to keep its value above a threshold (Carayannopoulos and Perez, 2015). Lehman Brothers Special Financing was the swap counterparty for four CAT bonds: Ajax Re, Carillon, Newton Re and Willow Re. When Lehman failed in September 2008, these bonds were downgraded and their collateral lost value. Ajax Re paid all its interest but could not repay its principal in full at maturity in May 2009 (Artemis, 2009). After issuance froze in late 2008, new deals restricted collateral to cash, government-guaranteed bank debt, U.S.

Treasury money-market funds and putable notes from sovereign or quasi-governmental issuers, with frequent mark-to-market and trustee monitoring (Carayannopoulos and Perez, 2015). No new CAT bond used a total return swap after 2009, and by the end of 2011 Treasury money-market funds accounted for more than 60 percent of collateral structures (Artemis, 2011).

Another is **liquidity**. CAT bonds trade over the counter and not very often. Using 9,393 secondary-market trades reported to TRACE between 2015 and 2019, Herrmann and Hibbeln (2023) find that liquidity is lower for riskier bonds and for bonds far from maturity, and they estimate that the liquidity premium accounts for about a fifth of the observed CAT bond yield spread, 98 basis points on average. A premium of that size can move when dealers and investors run short of capital, which is also when stock prices fall (Brunnermeier and Pedersen, 2009). In late 2008, distressed CAT bonds were put up for sale quickly in a market drained of liquidity, and investors suffered losses "not from natural catastrophic events, but from market risks that they thought they had been isolated from" (Carayannopoulos and Perez, 2015, pp. 22-23).

A third is **measurement**. The three articles use two Swiss Re indices, which are built from secondary-market bid indications rather than trades (Drobetz et al., 2020), and the net asset value of a CAT bond fund. An index of quotes can be smoothed, which lowers its measured volatility and its correlation with other markets (Getmansky, Lo and Makarov, 2004).

2.2 Diversifier, hedge and safe haven: what each test measures

The articles use, explicitly or not, the definitions of Baur and Lucey (2010). A hedge is uncorrelated or negatively correlated with another asset or portfolio on average. A diversifier is positively, but not perfectly, correlated with it on average. A safe haven is uncorrelated or negatively correlated with it during extreme market declines. Baur and McDermott (2010) add that a strong hedge or safe haven is negatively correlated with the market and a weak one is uncorrelated. The zero-beta claim, which says the beta is zero at all times, therefore amounts to a weak hedge plus a weak safe haven, and evidence of diversification does not confirm it. The definitions are also about correlation, not price: a statistical safe haven can lose money in a crash as long as its losses do not move with the market's.

Two warnings follow for reading the articles. The word "strong" changes meaning between studies: in the regression test used by Drobetz et al. (2020), a strong safe haven is an asset whose correlation falls significantly in extreme weeks, which does not require that correlation to be negative. And the tools differ: a beta divides the covariance with the market by the market's variance, a correlation scales the same covariance by both volatilities, and a copula describes dependence separately from the individual distributions, including in the tails. They can give different verdicts on the same data (Section 3.4).

2.3 Carayannopoulos and Perez (2015): zero beta, except in a crisis

The first article asks whether CAT bonds are zero-beta investments and whether they diversify a portfolio. The authors use 614 weekly returns, from January 2002 to October 2013, on the Swiss Re BB CAT bond price index, the S&P 500, an investment-grade U.S. corporate bond index and

a 5-to-7-year government bond index. They estimate time-varying covariances with a diagonal-VEC multivariate GARCH model, identify high- and low-correlation regimes with a Markov-switching model, and compute dynamic hedge ratios, the conditional covariance with the S&P 500 divided by its conditional variance, which they read as time-varying betas. The crisis is the NBER period from December 2007 to May 2009.

Simple correlations already show the pattern. The correlation between CAT bonds and the S&P 500 is 0.016 before the crisis, 0.252 during it (significant at 5 percent) and 0.058 after it; with corporate bonds it goes from 0.003 to 0.475 and back to 0.092. With government bonds, by contrast, it falls from 0.128 to an insignificant 0.054: in the crisis, CAT bonds moved with credit, not with safe fixed income. The dynamic estimates add the timing. The correlation with the stock market stayed close to zero until the end of August 2008, rose from September 2008 to 0.24 in May 2009, fluctuated widely through 2010, and returned to its pre-crisis level at the beginning of 2011. The high-correlation regime that began in October 2008 lasted about 30 months; those after Katrina, the 2011 Japanese earthquake and Sandy were much shorter. The economic size of the effect was small. The CAT bond beta peaked at 0.014 in January 2010, against 0.13 for corporate bonds, and the weekly volatility of CAT bonds barely moved (0.37 percent before the crisis, 0.41 percent during it) while that of the S&P 500 rose from 1.91 to 4.49 percent.

The authors conclude that CAT bonds "are not zero-beta assets" (p. 21), that weaknesses of the trust accounts were the main cause, and that CAT bonds remain a valuable diversifier because the crisis affected them far less than other assets. Their results hold with the broad Swiss Re index and with the BB total return index, and they describe the effect of the crisis as temporary. They also note that Hurricane Ike struck Texas in the same month as Lehman's failure and argue, by comparison with the other events, that Ike acted as a catalyst but cannot explain the size and persistence of the 2008 correlation.

2.4 Drobetz, Schröder and Tegtmeier (2020): a diversifier, and a safe haven only after the crisis

The second article widens the comparison to an international multi-asset portfolio. The authors pair the Swiss Re Global Unhedged CAT Bond Performance Index with total return indices for world stocks (MSCI World IMI), global government bonds, listed real estate, commodities, listed private equity and infrastructure. They use 885 weekly returns from January 2002 to December 2018, split into a pre-crisis period (2002-2008) and a post-crisis period (2009-2018). Following Bouri et al. (2017), they first estimate dynamic conditional correlations (DCC) between CAT bonds and each asset class, then regress them on dummies for the weeks in which the other asset's return falls below its 10, 5 and 1 percent quantiles: $DCC_t = m_0 + m_1 \cdot d_{10\%} + m_2 \cdot d_{5\%} + m_3 \cdot d_{1\%}$. A significantly positive m_0 makes CAT bonds a diversifier, a zero (negative) m_0 a weak (strong) hedge. Insignificant dummy coefficients make them a weak safe haven, significantly negative ones a strong safe haven.

The constant m_0 is positive and significant at 1 percent for every asset class: about 0.11 with stocks over the full period, 0.24 before 2009 and 0.10 after. CAT bonds are therefore a diversifier

against all six asset classes and a hedge against none. Before 2009 several dummy coefficients are positive, including 0.072 for the 5 percent stock quantile, so the correlation rose in extreme weeks. After 2009 only one coefficient is negative and significant: m_2 for stocks, −0.091 at the 10 percent level, or −0.101 at the 5 percent level with desmoothed returns. On this coefficient rests the conclusion that CAT bonds are a strong safe haven against extreme stock market declines only in the post-crisis period. The authors also conclude that CAT bonds are "not immune to changes in systematic risk" and that their behaviour "is not consistent with a zero-beta investment" (p. 5).

The summary statistics show the two faces of the asset. The CAT bond index has the lowest weekly volatility of the seven series (0.76 percent) but the highest kurtosis (313.5) and a strongly negative skewness (−9.0), the mark of rare but large event losses. Its first-order autocorrelation is significantly negative, which the authors attribute to the indicative pricing of the index; they re-estimate every regression on desmoothed returns and find the estimates "virtually unchanged" (p. 5).

2.5 Haffar and Le Fur (2022): little extreme co-movement

The third article studies the dependence structure. Its CAT bond series is not an index but a CAT bond fund, the GAMSTAR Institutional USD Acc share class managed by GAM, whose monthly returns are compared with six stock indices (CAC 40, DJIA, EURO STOXX 50, FTSE 100, Hang Seng, Nikkei 225) over 100 months, from November 2011 to February 2020. The authors model each series with an ARMA-GARCH process and join them with a Student-t copula with asymmetric dynamic conditional correlation. They then compare risk-adjusted performance, contributions to the portfolio's value-at-risk and expected shortfall, co-moments with an equally weighted benchmark portfolio of the other assets, and mean-variance optimal portfolios. They present the paper as the first to use copulas for CAT bonds.

Their findings are the most favourable of the three. The fund has the highest modified Sharpe ratio (0.112, against 0.027 to 0.089 for the stock indices). Its beta with the benchmark portfolio is 0.031, against 0.77 to 1.27 for the stock indices, and its "beta co-kurtosis" is 0.015, against 0.80 to 1.35. Its contribution to the portfolio's 99 percent value-at-risk is negative (−0.54 percent). From the co-kurtosis result, the authors conclude that CAT bonds are the asset that raises the probability of extreme co-variations with the benchmark portfolio the least. They cite Carayannopoulos and Perez (2015) in support (p. 298) and conclude that their results "confirm those of previous studies" (p. 307).

2.6 Other studies: divergent and supporting evidence

Two studies of CAT bond prices take the opposite side. Gürtler et al. (2016) find that the financial crisis and Hurricane Katrina both had a significant effect on secondary-market premiums, and that premiums are positively related to corporate bond spreads, a link that grew stronger during the crisis; they conclude that CAT bonds should not be regarded as zero-beta securities. Braun (2016) finds that the spreads of comparably rated corporate bonds help explain

CAT bond spreads at issuance. Both show that the price investors demand for catastrophe risk moves with the price of credit risk.

A recent study supports the favourable view for a later crisis. Applying a DCC-GARCH model to a multi-asset portfolio before and during the COVID-19 pandemic, Feng and Zeng (2025) find that CAT bonds were a diversifier over the whole sample and a safe haven against stocks and private equity after the outbreak. The results hold with desmoothed returns, but hedge ratios are sensitive to hurricane risk. The same pandemic produced a case in which the zero-beta argument fails by design: the World Bank's 2017 pandemic bonds put 16.67 percent of the principal of the Class A notes and all of the principal of the Class B notes at risk if a qualifying pandemic occurred (Government Actuary's Department, 2021), and COVID-19 met the trigger conditions during the same shock that hit equity markets. Table 1 summarises the three main articles.

Table 1. The three articles compared

	Carayannopoulos and Perez (2015)	**Drobetz, Schröder and Tegtmeier (2020)**	**Haffar and Le Fur (2022)**
Question	Are CAT bonds zero-beta, and did the subprime crisis change this?	Are CAT bonds a diversifier, a hedge or a safe haven?	How do CAT bonds depend on equity markets, including in the extremes?
Data	Swiss Re BB price index vs S&P 500, corporate and government bonds; 614 weekly returns, 2002-2013	Swiss Re Global Unhedged index vs six asset classes; 885 weekly returns, 2002-2018	GAMSTAR CAT bond fund vs six stock indices; 100 monthly returns, Nov. 2011 to Feb. 2020
Method	Diagonal-VEC MGARCH, Markov switching, dynamic hedge ratios	DCC-GARCH, then regression of the DCC on 10/5/1 percent quantile dummies	Student-t copula with asymmetric DCC; performance, VaR, co-moments
Stress	NBER crisis, Dec. 2007 to May 2009	Extreme weeks of the other asset; before vs after 2009	Co-kurtosis; no crisis in the sample
Key results	Correlation with S&P 500 before / during / after: 0.016 / 0.252 / 0.058; beta peaks at 0.014	Average correlation with stocks 0.24 before 2009, 0.10 after; post-crisis $m_2 = -0.091$	Beta 0.031; beta co-kurtosis 0.015; modified Sharpe ratio 0.112
Position on zero beta	Rejected in the crisis, holds outside it	Rejected on average; safe haven against stocks after 2009	Closest to zero beta (not tested formally)

Source: compiled by the author from the three articles.

3. Analysis

3.1 How far the evidence of each article goes

Each article answers its own question correctly, but none of them answers the question of this paper on its own, and the limits differ.

The strength of Carayannopoulos and Perez is that they treat beta as a variable and date the change precisely: the correlation stayed near zero through the first nine months of the recession and jumped only when Lehman's failure hit the swap structures. Their comparison of bond indices also locates the shock in credit. Their limit is that a beta estimated on returns cannot say

which part of the bond moved with the market. A fall in the value of Lehman-guaranteed collateral and a rise in the premium demanded by investors who needed cash both appear as a positive beta, and the authors' own account of 2008 describes both. Their treatment of Hurricane Ike is reasonable, but it is an argument by comparison, not a formal separation of the two shocks.

Drobetz et al. come closest to my question, because they separate the three properties and test them by sub-period. Three features limit their safe-haven result. It rests on one coefficient out of three. Because the dummies are nested, the implied correlation in the worst 5 percent of post-crisis weeks is $m_0 + m_1 + m_2 \approx 0.10 + 0.04 - 0.09 \approx 0.05$, or about 0.01 if the insignificant m_1 is set to zero, and the insignificant m_3 is positive (0.10). Either way, this is at most a weak safe haven in Baur and McDermott's terms, not a strong one. Second, the pre-crisis period ends in December 2008, so it contains the months after Lehman's failure. Third, the dummies identify stress through the other asset's return only, so a week of falling stock prices with working funding markets and a week of forced selling are treated alike.

Haffar and Le Fur's favourable result has the narrowest base. The extreme co-movement finding comes from co-kurtosis betas, not from a copula tail-dependence coefficient, and the copula's correlation dynamics are weakly identified: the DCC parameter a is estimated at zero, which the authors note is inconsistent with their test of dynamic correlation. The performance advantage is not statistically distinct, since their peer-performance test gives a 98 percent probability that the fund performs equally to the stock indices, and their most-diversified portfolio gives the fund a zero weight. Ranking first among six stock indices, which depend strongly on each other and on an equity benchmark by construction, is a modest test. The series is a fund that holds 92 percent fixed income and 8 percent cash, so its returns mix CAT bond risk with the manager's choices, and the paper dates the fund's launch to 26 March 2012, four months after its return series begins, without explaining the earlier prices. The authors themselves read their beta cautiously: there is "a slight positive beta co-variance", so adding the fund to the benchmark portfolio "will not reduce its volatility" (pp. 304-305). Most important for my question, the sample ends on 28 February 2020, days before the March 2020 sell-off, and contains no funding crisis and no collateral failure.

Once these limits are taken into account, the contradictions are of two kinds. The first concerns average behaviour. Outside crises, Carayannopoulos and Perez cannot distinguish the correlation of CAT bonds with the stock market from zero, a weak hedge in Baur and McDermott's terms, while Drobetz et al. find a positive and significant average correlation, a diversifier but not a hedge. The second and sharper contradiction concerns stress: CAT bonds moved with markets after Lehman's failure and in extreme weeks before 2009, but Drobetz et al. find a safe haven after 2009, Feng and Zeng after the COVID-19 outbreak, and Haffar and Le Fur the weakest extreme co-movement of all the assets they study. Haffar and Le Fur see no contradiction, because their question is diversification; the disagreement appears only when the question is the zero-beta property.

The numbers show that the disagreement is smaller than the labels suggest. The reported betas are small, 0.014 at the 2010 peak and 0.031 in calm years, although they come from different

series. The studies disagree mainly on whether a small but significant exposure contradicts the zero-beta claim, and on whether 2008 tells us anything about later crises. A simple account would call 2008 a special case and the rest a matter of method. My argument is that the bonds changed, the crises differed and the methods differed, and that a decomposition separates the first two and makes a prediction the simple account does not (Section 3.5).

3.2 My contribution: a decomposition of the CAT bond return

I write the return of a CAT bond, or of an index of CAT bonds, over one period as

$$R_t = r_t + s - \ell_t - D \cdot \Delta\pi_t - k_t \tag{1}$$

where r_t is the return on the collateral, s is the contractual spread accrued over the period, ℓ_t is the loss caused by catastrophes (a realised loss, or a revaluation of the expected loss after an event), $\Delta\pi_t$ is the change in the premium π that investors require above the expected loss, D is the bond's spread duration, and k_t is the loss in collateral value caused by the failure of a counterparty or a fall in the value of the collateral assets. Defining π net of the expected loss keeps a revision of the expected loss in ℓ_t only. The insurance core is $s - \ell_t$ and the financial wrapper is $r_t - D \cdot \Delta\pi_t - k_t$.

The beta of the bond with respect to the market return R_m is the sum of the betas of its components, $\beta_x = \mathrm{Cov}(x_t, R_{m,t})/\mathrm{Var}(R_{m,t})$, where β_π is the beta of $\Delta\pi_t$:

$$\beta = \beta_r - \beta_\ell - D \cdot \beta_\pi - \beta_k \tag{2}$$

The zero-beta argument is an assumption about ℓ_t only: natural catastrophes do not depend on financial markets, so $\beta_\ell = 0$. The collateral earns a short-term money-market rate whose weekly variation is negligible, so $\beta_r \approx 0$, and s is fixed by contract. What is left is

$$\beta \approx -D \cdot \beta_\pi - \beta_k \tag{3}$$

Under this assumption, the measured beta of a natural-catastrophe bond is the beta of its financial wrapper. I call the first term the premium channel and the second the collateral channel. When markets fall, the required premium rises ($\beta_\pi < 0$) and collateral losses, if possible at all, become more likely ($\beta_k < 0$), so both channels push the beta above zero in a crisis. The two channels differ in what they cost the investor. As in Campbell and Vuolteenaho's (2004) split between cash-flow and discount-rate beta, a price fall caused by a higher premium is recovered as the bond approaches maturity if no catastrophe occurs, so it is a loss mainly for a holder who has to sell; a collateral loss is permanent.

Two qualifications apply. Carayannopoulos and Perez find short high-correlation regimes after Katrina, the 2011 Japanese earthquake and Sandy, so β_ℓ is small rather than exactly zero. And the premium channel is not free of catastrophe risk in every sense, since π is the price investors charge for bearing it. What the decomposition separates is the occurrence of catastrophes, which markets do not drive, from the price and the packaging of catastrophe risk, which they can.

Equation (3) leads to three predictions.

1. In calm markets the collateral channel is closed and the premium moves little, so β is small. It need not be exactly zero: the link between CAT premiums and corporate spreads found by Gürtler et al. (2016) and Braun (2016) exists outside crises, which is consistent with the average correlation of about 0.10 found by Drobetz et al. after 2009.
2. The collateral channel can open only when the collateral is exposed to the credit of a financial institution. It could open before 2009 and is close to closed for deals whose collateral is held in Treasury money-market funds or government-backed notes.
3. The premium channel opens when holders of CAT bonds need cash at the moment the market falls. They sell, and dealers short of capital absorb the bonds only at a higher premium. When funding is scarce, market liquidity and funding liquidity reinforce each other (Brunnermeier and Pedersen, 2009), and when capital moves slowly into a market after such a shock, prices fall further than fundamentals justify and then recover gradually (Duffie, 2010). If the capital available to bear catastrophe risk shrinks with financial wealth, its price also rises with credit spreads, as Froot's (2001) account of capital-driven pricing predicts.

The decomposition also identifies the case in which the zero-beta claim fails at its root. If the peril is itself a macroeconomic shock, $\beta_\ell \neq 0$. For diversification purposes, pandemic bonds, and in my view bonds covering systemic cyber risk, belong to a different asset class from natural-catastrophe bonds.

3.3 Four stress regimes

Equation (3) suggests classifying stress episodes on two dimensions: whether the collateral was exposed to bank credit, and whether the episode was a funding crisis, in which investors needed cash, or a fall in equity prices without forced selling. Table 2 places the evidence in this grid.

Table 2. Stress regimes implied by the decomposition, with the evidence for each

	Funding crisis (investors need cash)	**Equity sell-off, no forced selling**
Collateral exposed to bank credit (swap structures issued before 2009)	Both channels open: beta turns positive. *Evidence:* correlation with the S&P 500 jumps after September 2008 (Carayannopoulos and Perez, 2015); correlation rises in extreme weeks before 2009 (Drobetz et al., 2020).	Only the insurance core is tested: beta close to zero. *Evidence:* correlation near zero through the first nine months of the recession (Carayannopoulos and Perez, 2015).
Collateral in Treasury funds or government-backed notes (deals issued after 2009)	Premium channel only: a small fall, later reversed; at most a weak safe haven. *Evidence:* safe haven against stocks after the COVID-19 outbreak (Feng and Zeng, 2025); fall of 0.81 percent in March 2020 (Insurance Insider ILS, 2020).	Neither channel open: beta close to zero, weak safe haven. *Evidence:* implied correlation of 0.01 to 0.05 in the worst post-2009 stock weeks (Drobetz et al., 2020); beta of 0.031 over 2011-2020 (Haffar and Le Fur, 2022).

Source: classification by the author, based on the studies cited in each cell.

Each of the main results can be placed in one cell. Since the grid was built after reading the evidence, this fit shows compatibility, not proof.

The top-left cell contains one episode, the months after Lehman's failure. The timing found by Carayannopoulos and Perez is consistent with the collateral channel but cannot isolate it: no forced selling of CAT bonds is documented in the first nine months of the recession, when the correlation stayed near zero, and the correlation jumped only when the swap structures failed. The positive extreme-week coefficients of Drobetz et al. before 2009 belong to the same cell, since their pre-crisis period ends in December 2008. The persistence of the high-correlation regime until early 2011 also fits: swap-backed bonds stayed outstanding until maturity, and capital returned slowly (Duffie, 2010).

The bottom-right cell explains the favourable results. As swap-backed deals matured after 2009, the collateral channel closed, and I assume, as the test of Section 3.5 can check, that post-crisis stock declines brought no forced selling on the scale of late 2008, so the premium channel stayed weak. The grid predicts a beta close to zero and a weak safe haven, which is what the numbers show. The whole sample of Haffar and Le Fur lies in this cell. Their favourable result is what the grid predicts for it, and for the same reason it says nothing about funding crises.

March 2020 is the main test of the bottom-left cell: collateral in Treasury funds combined with a severe scramble for cash. The framework predicts a fall driven by forced sales, no loss on collateral, and a recovery once capital returns. Market reports describe such a fall: the Swiss Re price return index fell by 0.81 percent in the week to 20 March, with markdowns linked to bulk lists of bonds put up for sale (Insurance Insider ILS, 2020), and by May the sell-off was described as temporary (The Insurer, 2020). A bond with a maturity of three to four years has a spread duration of roughly one to two years. If the whole fall came from the premium, equation (1) turns it into a rise of about 40 to 80 basis points (0.81 percent divided by a duration of two and one years), of the same order as the liquidity premium of 98 basis points estimated by Herrmann and Hibbeln (2023). On this reading the premium channel opened in March 2020 but stayed small, which is consistent with Feng and Zeng still finding a safe haven. Their result fits the simple account as well as mine; what the simple account does not predict is a fall linked to forced sales and then reversed, and that is the part of the evidence the decomposition explains.

3.4 Measurement choices that move the answer

Part of the divergence between the articles is statistical. The first source is the gap between economic and statistical significance. CAT bond returns are much less volatile than stock returns, so a significant correlation can go with a tiny beta. In the data of Carayannopoulos and Perez, the crisis correlation of 0.25 goes with a beta that peaked at 0.014: a 20 percent fall in the S&P 500 would have cost a CAT bond investor about 0.3 percent. The zero-beta claim is rejected statistically and close to true economically, and the articles differ partly in which of the two they stress.

The second is conditioning on volatility. Suppose the CAT bond return is β times the market return plus an independent shock with volatility σ_ε. The correlation is then $\rho = \beta \cdot \sigma_m / \sqrt{(\beta^2 \sigma_m^2 + \sigma_\varepsilon^2)}$, which rises with the market's volatility σ_m whenever $\beta > 0$ (Forbes and Rigobon, 2002). In the data of Carayannopoulos and Perez, the weekly volatility of the S&P 500 rose from 1.91 to 4.49

percent, which at a constant beta and idiosyncratic volatility would lift the correlation only from 0.016 to under 0.04. The rest of the rise to 0.252 is not mechanical, and their beta, which is free of this bias, also rose.

The third is smoothing, which matters less than theory suggests. A smoothed index biases beta and correlation toward zero (Getmansky et al., 2004), but the estimates of Drobetz et al. and of Feng and Zeng hold with desmoothed returns. Drobetz et al. attribute the negative first-order autocorrelation of their index (−0.34) to smoothing, whereas smoothing produces positive autocorrelation in the model of Getmansky et al. A negative autocorrelation is the mark of noise in quoted prices that reverses the following week (Roll, 1984). Such noise adds variance instead of removing it, so the bias it creates is not the one that desmoothing corrects.

The fourth is the choice of proxy, measure and window, and it matters most. Before the crisis, the correlation with stocks is 0.016 for the BB price index and the S&P 500 in Carayannopoulos and Perez, but about 0.24 for the global index and world stocks in Drobetz et al. The few weeks of late 2008 in the second sample cannot explain a gap of that size. It has to come from the series (a BB-rated price index against a global total return index), the measure (a sample correlation against the mean of a DCC series) and the comparison index. A monthly fund series over a calm window, finally, cannot detect what a weekly index containing 2008 detects. These differences explain why the three studies can all be right on their own data.

3.5 Proposed solutions

Tests that could prove the framework wrong. No estimation is done here, but the framework specifies tests that would settle the disagreement rather than add one more estimate to it.

- ***Split the 2008 beta between the two channels.*** Compare the four Lehman-swap bonds with other CAT bonds from September to December 2008. A gap between the two groups measures the collateral channel, plus any extra discount on distressed bonds. A fall common to all bonds, larger for bonds with a longer spread duration, measures the premium channel.
- ***Measure funding stress directly.*** Estimate $DCC_t = c_0 + c_1 \cdot Ext_t + c_2 \cdot Fund_t + c_3 \cdot (Ext_t \times Fund_t)$, where Ext_t marks the worst 5 percent of stock weeks and $Fund_t$ marks the top decile of funding stress, measured for example by the spread between commercial paper and Treasury bills. My framework predicts c_1 close to zero and $c_3 > 0$ in both periods, smaller after 2009, when only the premium channel can open. The simple account, in which 2008 was a special case, predicts $c_3 = 0$ after 2009. March 2020 and September-October 2022, when the UK pension funds' liability-driven investment crisis coincided with Hurricane Ian, provide the observations.
- ***Re-run the favourable design on a sample with a funding crisis.*** Estimating the copula and co-kurtosis measures of Haffar and Le Fur on weekly index data that include 2008 and March 2020 would test the grid's prediction of stronger extreme co-movement in funding-crisis weeks.

- ***Report betas next to correlations,*** adjust crisis correlations for volatility (Forbes and Rigobon, 2002), use TRACE transaction prices where possible (Herrmann and Hibbeln, 2023), and enter weeks with a major catastrophe, such as Ike in 2008 and Ian in 2022, as separate dummies, so that a coincidence of shocks is not read as market beta.

Portfolio and market rules that remove what can be removed.

- ***Keep the collateral channel closed.*** Investors should require collateral in Treasury money-market funds or government-backed notes, held in trust, marked to market and disclosed. Most of the market has moved this way since 2009 (Artemis, 2011), but investors should verify the terms of each issue.
- ***Contain the premium channel.*** An investor that plans to raise cash by selling CAT bonds in a crisis turns a temporary markdown into a realised loss. Funds that promise frequent redemptions should match those terms to the liquidity of their bonds through notice periods, cash buffers or redemption limits. Risk managers who size the allocation with value-at-risk should stress it on a funding crisis such as March 2020, not only on calm data.
- ***Limit concentration and classify perils honestly.*** Limits by peril, region and sponsor keep the insurance core diversified, which matters since Feng and Zeng (2025) find hedge ratios sensitive to hurricane risk. Bonds on macroeconomic perils, such as pandemics, should not be presented or modelled as zero-beta exposures.

3.6 Conclusion of the analysis

Whether CAT bonds are zero-beta assets depends on which part of the bond the question refers to. If natural catastrophes do not depend on markets, the insurance core is zero-beta. The traded bond has a small beta in calm markets, a beta that can rise in a funding crisis through the premium channel, and a beta that rose in 2008, when both channels opened, although even then the exposure was economically small. For the literature, this places the link between CAT premiums and corporate spreads in the financial wrapper, not in the catastrophe losses, and makes the safe-haven property depend on the structure of the asset and the type of stress, not only on the state of the market. It also makes CAT bonds a rare market in which a liquidity-driven fall can in principle be told apart from a fundamental loss, because the catastrophes that cause fundamental losses are observable.

4. Conclusion

This paper asked to what extent CAT bonds retain their diversification benefits during financial stress, and why studies disagree. The three main articles reach contradictory conclusions on the zero-beta property and on crisis behaviour. My contribution is a framework that explains why: it splits the CAT bond return into an insurance core, close to independent of markets for natural perils, and a financial wrapper, whose exposure depends on the collateral structure and on whether a stress episode is a funding crisis. The three articles sample different cells of Table 2

with different proxies, frequencies and windows, so their conclusions can diverge while their estimates are compatible.

The answer to the research question is therefore conditional. CAT bonds kept their diversification benefit in every episode reviewed: even after Lehman's failure, their beta peaked at 0.014. Their safe-haven property holds, weakly, against stocks when the collateral is safe and holders are not forced to sell, and the framework predicts that it erodes when either condition fails. Investors can make the benefit more reliable through collateral rules, liquidity management and concentration limits, but they should treat CAT bonds as a diversifier with a conditional safe-haven property, not as crisis insurance.

The analysis has limits. It rests on published results rather than new estimates, and the grid was built after reading the evidence, so the tests of Section 3.5 still need to be run. The bottom-left cell of Table 2 rests mainly on one episode, March 2020. Finally, the framework assumes that catastrophe losses are nearly independent of markets for natural perils, an assumption that holds only approximately and does not extend to pandemic or systemic cyber risks.

Bibliography

Artemis (2009). "Ajax Re notes downgraded again by S&P after principal not paid in full," *Artemis.bm*, May 12, 2009. https://www.artemis.bm/news/ajax-re-notes-downgraded-again-by-sp-after-principal-not-paid-in-full/

Artemis (2011). "The death of the total return swap and the rise of the Treasury money market fund," *Artemis.bm*, December 19, 2011. https://www.artemis.bm/news/the-death-of-the-total-return-swap-and-the-rise-of-the-treasury-money-market-fund/

Baur, Dirk G., and Brian M. Lucey (2010). "Is gold a hedge or a safe haven? An analysis of stocks, bonds and gold," *The Financial Review*, vol. 45, no. 2, pp. 217–229. https://doi.org/10.1111/j.1540-6288.2010.00244.x

Baur, Dirk G., and Thomas K. McDermott (2010). "Is gold a safe haven? International evidence," *Journal of Banking and Finance*, vol. 34, no. 8, pp. 1886–1898. https://doi.org/10.1016/j.jbankfin.2009.12.008

Bouri, Elie, Peter Molnár, Georges Azzi, David Roubaud, and Lars Ivar Hagfors (2017). "On the hedge and safe haven properties of Bitcoin: Is it really more than a diversifier?" *Finance Research Letters*, vol. 20, pp. 192–198. https://doi.org/10.1016/j.frl.2016.09.025

Braun, Alexander (2016). "Pricing in the primary market for cat bonds: New empirical evidence," *Journal of Risk and Insurance*, vol. 83, no. 4, pp. 811–847. https://doi.org/10.1111/jori.12067

Brunnermeier, Markus K., and Lasse Heje Pedersen (2009). "Market liquidity and funding liquidity," *Review of Financial Studies*, vol. 22, no. 6, pp. 2201–2238. https://doi.org/10.1093/rfs/hhn098

Campbell, John Y., and Tuomo Vuolteenaho (2004). "Bad beta, good beta," *American Economic Review*, vol. 94, no. 5, pp. 1249–1275. https://doi.org/10.1257/0002828043052240

Carayannopoulos, Peter, and M. Fabricio Perez (2015). "Diversification through catastrophe bonds: Lessons from the subprime financial crisis," *The Geneva Papers on Risk and Insurance – Issues and Practice*, vol. 40, no. 1, pp. 1–28. https://doi.org/10.1057/gpp.2014.14

Cummins, J. David (2008). "CAT bonds and other risk-linked securities: State of the market and recent developments," *Risk Management and Insurance Review*, vol. 11, no. 1, pp. 23–47. https://doi.org/10.1111/j.1540-6296.2008.00127.x

Cummins, J. David, and Mary A. Weiss (2009). “Convergence of insurance and financial markets: Hybrid and securitized risk-transfer solutions,” *Journal of Risk and Insurance*, vol. 76, no. 3, pp. 493–545. https://doi.org/10.1111/j.1539-6975.2009.01311.x

Drobetz, Wolfgang, Henning Schröder, and Lars Tegtmeier (2020). “The role of catastrophe bonds in an international multi-asset portfolio: Diversifier, hedge, or safe haven?” *Finance Research Letters*, vol. 33, article 101198. https://doi.org/10.1016/j.frl.2019.05.016

Duffie, Darrell (2010). “Presidential address: Asset price dynamics with slow-moving capital,” *Journal of Finance*, vol. 65, no. 4, pp. 1237–1267. https://doi.org/10.1111/j.1540-6261.2010.01569.x

Feng, Chi, and Xudong Zeng (2025). “The Covid-19 pandemic and the portfolio diversification effect of catastrophe bonds,” *Journal of Risk*, vol. 27, no. 4, pp. 85–112. https://doi.org/10.21314/JOR.2024.022

Forbes, Kristin J., and Roberto Rigobon (2002). “No contagion, only interdependence: Measuring stock market comovements,” *Journal of Finance*, vol. 57, no. 5, pp. 2223–2261. https://doi.org/10.1111/0022-1082.00494

Froot, Kenneth A. (2001). “The market for catastrophe risk: A clinical examination,” *Journal of Financial Economics*, vol. 60, nos. 2–3, pp. 529–571. https://doi.org/10.1016/S0304-405X(01)00052-6

Getmansky, Mila, Andrew W. Lo, and Igor Makarov (2004). “An econometric model of serial correlation and illiquidity in hedge fund returns,” *Journal of Financial Economics*, vol. 74, no. 3, pp. 529–609. https://doi.org/10.1016/j.jfineco.2004.04.001

Government Actuary’s Department (2021). “Pandemic bonds: What are they and how do they work?” *GAD blog*, March 2, 2021. https://actuaries.blog.gov.uk/2021/03/02/pandemic-bonds-what-are-they-and-how-do-they-work/

Gürtler, Marc, Martin Hibbeln, and Christine Winkelvos (2016). “The impact of the financial crisis and natural catastrophes on CAT bonds,” *Journal of Risk and Insurance*, vol. 83, no. 3, pp. 579–612. https://doi.org/10.1111/jori.12057

Haffar, Adlane, and Éric Le Fur (2022). “Dependence structure of CAT bonds and portfolio diversification: A copula-GARCH approach,” *Journal of Asset Management*, vol. 23, no. 4, pp. 297–309. https://doi.org/10.1057/s41260-022-00271-3

Herrmann, Markus, and Martin Hibbeln (2023). “Trading and liquidity in the catastrophe bond market,” *Journal of Risk and Insurance*, vol. 90, no. 2, pp. 283–328. https://doi.org/10.1111/jori.12407

Insurance Insider ILS (2020). “Cat bond prices limit Covid-19 fall to under 1%,” *Insurance Insider ILS*, March 2020. https://www.insuranceinsiderils.com/article/2876fhy9ajjsuvp26rkuq/cat-bond-prices-limit-covid-19-fall-to-under-1

Roll, Richard (1984). “A simple implicit measure of the effective bid-ask spread in an efficient market,” *Journal of Finance*, vol. 39, no. 4, pp. 1127–1139. https://doi.org/10.1111/j.1540-6261.1984.tb03897.x

The Insurer (2020). “Cat bonds nudge up after temporary Covid-19 sell-off,” *The Insurer*, May 2020. https://theinsurer.com/analysis/cat-bonds-nudge-up-after-temporary-covid-19-sell-off